

OS HW1

資科工所 312551105 羅羽軒

- 文章位址: <https://hackmd.io/@lohsuan/rJlenpgg6>
- 題目敘述: [Compiling Linux Kernel and Adding Custom System Calls](#)

Prob 1: Change kernel suffix

Before compiling the kernel, you need to change your Linux kernel version suffix to `-os-<your-id>`, as an evidence that the kernel is built by yourself.

HW1-1 詳細過程: <https://hackmd.io/@lohsuan/Hk7C4rUWa>

Answer screenshot

```
yuhsuan@yuhsuan-virtual-machine: ~  
yuhsuan@yuhsuan-virtual-machine:~$ uname -a  
Linux yuhsuan-virtual-machine 5.19.12-os-312551105 #1 SMP PREEMPT_DYNAMIC Thu Oct 5 08:12:46 CST 2023 x86_64 x86_64 GNU/Linux  
yuhsuan@yuhsuan-virtual-machine:~$ cat /etc/os-release  
PRETTY_NAME="Ubuntu 22.04.3 LTS"  
NAME="Ubuntu"  
VERSION_ID="22.04"  
VERSION="22.04.3 LTS (Jammy Jellyfish)"  
VERSION_CODENAME=jammy  
ID=ubuntu  
ID_LIKE=debian  
HOME_URL="https://www.ubuntu.com/"  
SUPPORT_URL="https://help.ubuntu.com/"  
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"  
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"  
UBUNTU_CODENAME=jammy  
yuhsuan@yuhsuan-virtual-machine:~$ ^C  
yuhsuan@yuhsuan-virtual-machine:~$
```

Prob 2: Adding Custom System Calls

題目敘述: [Adding Custom System Calls](#)

In this part, you are required to implement two system calls on top of the kernel you built. One `sys_hello` and the other one `sys_revstr`.

2-1 sys_hello

1. Define new system call: sys_hello()

- mkdir workspace under `linux-5.15.12`,
- add custom system call file `hello.c` under `workspace` directory
 - `SYSCALL_DEFINE0`: 0 means no parameter pass in

```
// hello.c  
#include <linux/kernel.h>  
#include <linux/syscalls.h>  
SYSCALL_DEFINE0(hello)  
{  
    printk("Hello, world!\n");  
    printk("312551105\n");  
    return 0;  
}
```

- Create new `Makefile` in `workspace`. It will
 - specifies the objects to be built
 - added to the source during the next kernel recompilation.

```
// Makefile  
obj-y += hello.o
```

This is to ensure that the `hello.c` file is compiled and included in the kernel source code.

2. Add workspace/ to the kernel's Makefile

- switch back to `linux-5.19.12` directory
- add `workspace` directory to kernel's `Makefile`
 - find `core-y` line and append `workspace/`
 - this will tell the compiler to find our custom system call source code in `workspace/` directory

```
ifeq ($(KBUILD_EXTMOD),)  
core-y += kernel/certs/ mm/ fs/ ipc/ security/ crypto/  
core-$(CONFIG_BLOCK) += block/ workspace/  
core-$(CONFIG_IO_URING) += io_uring/
```

This is to tell the compiler that the source files of our new system call (`sys_hello()`) are in present in the `workspace` directory.

3. Add custom system call into the master's syscall table

- To wire up our new system call for x86 platforms, we need to update the master syscall tables.
 - Edit the system call table by `vim arch/x86/entry/syscalls/syscall_64.tbl`
 - add `sys_hello` with number 451

```
451    common    hello        sys_hello
```

4. Add function prototype in the system call header file

- The new entry point needs a corresponding function prototype
 - add it in `include/linux/syscalls.h`
 - marked as `asmlinkage` to match the way that system calls are invoked

```
asmlinkage long sys_hello(void);
```

5. Recompile the kernel

- Switch to the source directory `linux-5.19.12` and execute:

```
$ sudo make -j $(nproc)  
$ sudo make modules_install install
```

6. Reboot

- For the system to now use the newly configured kernel, `reboot`.

Test with TA's code

```
// hello_test.c  
#include <assert.h>  
#include <unistd.h>  
#include <sys/syscall.h>  
  
#define __NR_hello 451  
/* 451 is the system call number we declare in master's syscall table */  
int main(int argc, char *argv[]) {  
    int ret = syscall(__NR_hello);  
    assert(ret == 0);  
  
    return 0;  
}
```

Compile and run the program:

```
$ gcc hello_test.c  
$ ./a.out
```

show message with `sudo dmesg`

```
[ 517.054865] The reversed string: Hello, world!  
[ 517.054865] 312551105  
yuhsuan@yuhsuan-virtual-machine:~$
```

2-2 sys_revstr

The steps are the same as above

1. Define new system call: sys_revstr()

- add custom system call file `revstr.c` under `workspace` directory
 - `SYSCALL_DEFINE2`: 2 means 2 parameters pass in

```
// revstr.c  
#include <linux/kernel.h>  
#include <linux/syscalls.h>  
#include <linux/linkage.h>  
#include <linux/uaccess.h>  
  
SYSCALL_DEFINE2(revstr, int, len, char __user *, string)  
{  
    char str[200]; // declare the size of character string  
    unsigned long strlen = len;  
    char temp;  
  
    copy_from_user(str, string, strlen);  
    printk("The origin string: %s\n", str);  
  
    for (int i = 0; i < (strlen/2); i++) {  
        temp = str[i];  
        str[i] = str[strlen - i - 1];  
        str[strlen - i - 1] = temp;  
    }  
  
    printk("The reversed string: %s\n", str);  
  
    return 0;  
}
```

- Add `revstr.o` in `Makefile`

```
// Makefile  
obj-y += hello.o  
obj-y += revstr.o
```

2. Add custom system call into the master's syscall table

- To wire up our new system call for x86 platforms, we need to update the master syscall tables.
 - Edit the system call table by `vim arch/x86/entry/syscalls/syscall_64.tbl`
 - add `sys_revstr` with number 452

```
452    common    revstr        sys_revstr
```

3. Add function prototype in the system call header file

- The new entry point needs a corresponding function prototype
 - add it in `include/linux/syscalls.h`
 - marked as `asmlinkage` to match the way that system calls are invoked

```
asmlinkage long sys_revstr(int len, char __user *string);
```

4. Recompile the kernel

Switch to the source directory `linux-5.19.12` and execute:

```
$ sudo make -j $(nproc)  
$ sudo make modules_install install
```

5. Reboot

For the system to now use the newly configured kernel, `reboot`.

Test with TA's code

```
// revstr_test.c  
#include <assert.h>  
#include <unistd.h>  
#include <sys/syscall.h>  
  
#define __NR_revstr 452  
  
int main(int argc, char *argv[]) {  
    int ret1 = syscall(__NR_revstr, 5, "hello");  
    assert(ret1 == 0);  
  
    int ret2 = syscall(__NR_revstr, 11, "5Y573M C411");  
    assert(ret2 == 0);  
  
    return 0;  
}
```

Compile and run the program:

```
$ gcc revstr_test.c  
$ ./a.out
```

show message with `sudo dmesg`

```
[ 517.654800] The origin string: hello  
[ 517.654860] The reversed string: olleh  
[ 517.654864] The origin string: 5Y573M C411  
[ 517.654865] The reversed string: 114C M375Y5  
yuhsuan@yuhsuan-virtual-machine:~/kernelbuild/linux-5.19.12/workspace$
```